

20250419_01

April 19, 2025

```
[39]: import pandas as pd
import numpy as np

# Modeling and preprocessing
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline

# Evaluation
from sklearn.metrics import mean_absolute_error, root_mean_squared_error, r2_score

# Making dataset
df = pd.DataFrame({'student_id': [1, 2, 3, 4, 5],
                   'math_score': [75, 88, 95, 65, 50],
                   'english_score': [82, 79, 91, 70, 60],
                   'gender': ['F', 'M', 'M', 'F', 'F'],
                   'school_type': ['public', 'private', 'private', 'public', 'public'],
                   'final_score': [80, 85, 90, 70, 60]})
```

```
[41]: # Feature engineering
num_cols = ['math_score', 'english_score']
cat_cols = ['gender', 'school_type']

# Checking if any potential problems
{col: df[col].unique() for col in cat_cols}
```

```
[41]: {'gender': array(['F', 'M'], dtype=object),
      'school_type': array(['public', 'private'], dtype=object)}
```

```
[43]: # preprocessing
preprocessor = ColumnTransformer([('num', StandardScaler(), num_cols),
                                 ('cat', OneHotEncoder(drop = 'first'), cat_cols)])
```

1 First model : Linear Regression

```
[46]: # Creating pipeline
pipe_lr = Pipeline([('Preprocessing', preprocessor), ('model',
↳LinearRegression())])

# Making training and test sets
X = df.drop(columns = ['student_id', 'final_score'])
y = df['final_score']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.4,
↳random_state = 42)

# Training model
pipe_lr.fit(X_train, y_train)

# Prediction
y_pred_lr = pipe_lr.predict(X_test)
```

```
[48]: # Evaluating model
mae_lr = mean_absolute_error(y_test, y_pred_lr)
rmse_lr = root_mean_squared_error(y_test, y_pred_lr)
r2_lr = r2_score(y_test, y_pred_lr)

print('Linear Regression')
print('MAE: ', mae_lr)
print('RMSE: ', rmse_lr)
print('R²: ', r2_lr)
```

```
Linear Regression
MAE:  2.197117336581403
RMSE:  3.0376701200830367
R²:  0.9409443855459502
```

2 Second model : Decision Tree

```
[51]: # Creating another pipeline
pipe_tree = Pipeline([('Preprocessing', preprocessor),
↳('model', DecisionTreeRegressor(random_state = 42))])

# Training model
pipe_tree.fit(X_train, y_train)

# Prediction
y_pred_tree = pipe_tree.predict(X_test)
```

```
[53]: # Evaluating model
mae_tree = mean_absolute_error(y_test, y_pred_tree)
```

```
rmse_tree = root_mean_squared_error(y_test, y_pred_tree)
r2_tree = r2_score(y_test, y_pred_tree)

print('Decision Tree')
print('MAE: ', mae_tree)
print('RMSE: ', rmse_tree)
print('R2: ', r2_tree)
```

Decision Tree

MAE: 7.5

RMSE: 7.905694150420948

R²: 0.6

In this dataset, the linear regression model achieved significantly better performance than the decision tree regressor (MAE = 2.20 vs 7.50, R² = 0.94 vs 0.60). The tree-based model likely overfit due to the small dataset and limited feature diversity. For this problem, a linear model appears to be more appropriate.