

20250427_01

April 27, 2025

Today's objective : classify wines into 3 categories with Logistic Regression and Decision Tree.

```
[93]: # Loading the data and check how many are there
import pandas as pd
from sklearn.datasets import load_wine

data = load_wine()
X = pd.DataFrame(data.data, columns = data.feature_names)
y = pd.Series(data.target)

print(X.shape)
print(y.value_counts())
```

```
(178, 13)
```

```
1    71
```

```
0    59
```

```
2    48
```

```
Name: count, dtype: int64
```

X and y has the same number of rows, nice.

```
[95]: # Take a look at the dataset
X.describe()
```

```
[95]:
```

	alcohol	malic_acid	ash	alcalinity_of_ash	magnesium \
count	178.000000	178.000000	178.000000	178.000000	178.000000
mean	13.000618	2.336348	2.366517	19.494944	99.741573
std	0.811827	1.117146	0.274344	3.339564	14.282484
min	11.030000	0.740000	1.360000	10.600000	70.000000
25%	12.362500	1.602500	2.210000	17.200000	88.000000
50%	13.050000	1.865000	2.360000	19.500000	98.000000
75%	13.677500	3.082500	2.557500	21.500000	107.000000
max	14.830000	5.800000	3.230000	30.000000	162.000000

	total_phenols	flavanoids	nonflavanoid_phenols	proanthocyanins \
count	178.000000	178.000000	178.000000	178.000000
mean	2.295112	2.029270	0.361854	1.590899
std	0.625851	0.998859	0.124453	0.572359
min	0.980000	0.340000	0.130000	0.410000

25%	1.742500	1.205000	0.270000	1.250000
50%	2.355000	2.135000	0.340000	1.555000
75%	2.800000	2.875000	0.437500	1.950000
max	3.880000	5.080000	0.660000	3.580000

	color_intensity	hue	od280/od315_of_diluted_wines	proline
count	178.000000	178.000000	178.000000	178.000000
mean	5.058090	0.957449	2.611685	746.893258
std	2.318286	0.228572	0.709990	314.907474
min	1.280000	0.480000	1.270000	278.000000
25%	3.220000	0.782500	1.937500	500.500000
50%	4.690000	0.965000	2.780000	673.500000
75%	6.200000	1.120000	3.170000	985.000000
max	13.000000	1.710000	4.000000	1680.000000

X has no missing values, nice. But the value of each features varies a lot, need to standardize (for Logistic Regression).

```
[97]: # Making pipelines (Logistic Regression and Decision Tree)
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier

logreg_pipeline = Pipeline([('scaler', StandardScaler()), ('logreg',
↳ LogisticRegression(solver = 'liblinear', random_state = 42))])

tree_pipeline = Pipeline([('tree', DecisionTreeClassifier(random_state = 42))])
```

Now we train and test.

```
[99]: # Making training and testing sets
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2,
↳ random_state = 42)
```

```
[107]: # Training the models
from sklearn.model_selection import GridSearchCV

logreg_parm_grid = {'logreg__C': [0.01, 0.1, 1, 10, 100]}
logreg_grid = GridSearchCV(logreg_pipeline, logreg_parm_grid, cv = 5)
logreg_grid.fit(X_train, y_train)

tree_parm_grid = {'tree__max_depth': [3, 5, 10, None], 'tree__min_samples_leaf':
↳ [1, 3, 5]}
tree_grid = GridSearchCV(tree_pipeline, tree_parm_grid, cv = 5)
tree_grid.fit(X_train, y_train)
```

```
[107]: GridSearchCV(cv=5,
                  estimator=Pipeline(steps=[('tree',
DecisionTreeClassifier(random_state=42))]),
                  param_grid={'tree__max_depth': [3, 5, 10, None],
                              'tree__min_samples_leaf': [1, 3, 5]})
```

```
[128]: # Checking which one is the best and how best is it for each model
logreg_cv_results = pd.DataFrame(logreg_grid.cv_results_)
logreg_cv_results[['param_logreg__C', 'mean_test_score', 'std_test_score', 'rank_test_score']]
```

```
[128]:
```

	param_logreg__C	mean_test_score	std_test_score	rank_test_score
0	0.01	0.957635	0.026864	5
1	0.10	0.978818	0.017301	3
2	1.00	0.985961	0.017199	1
3	10.00	0.979064	0.017100	2
4	100.00	0.965025	0.021817	4

```
[117]: tree_cv_results = pd.DataFrame(tree_grid.cv_results_)
tree_cv_results[['param_tree__max_depth', 'param_tree__min_samples_leaf', 'mean_test_score', 'std_test_score', 'rank_test_score']]
```

```
[117]:
```

	param_tree__max_depth	param_tree__min_samples_leaf	mean_test_score	\
0	3	1	0.922414	
1	3	3	0.908374	
2	3	5	0.900985	
3	5	1	0.915271	
4	5	3	0.915517	
5	5	5	0.900985	
6	10	1	0.915271	
7	10	3	0.915517	
8	10	5	0.900985	
9	None	1	0.915271	
10	None	3	0.915517	
11	None	5	0.900985	

	std_test_score	rank_test_score
0	0.014819	1
1	0.017566	8
2	0.042337	9
3	0.018323	5
4	0.017057	2
5	0.042337	9
6	0.018323	5
7	0.017057	2
8	0.042337	9
9	0.018323	5

10	0.017057	2
11	0.042337	9

Since the best params for Logistic Regression(0.985961) outperform the best of Decision Tree(0.922414), I decide to use Logisitic Regression for the final test set.

```
[134]: # Predicting and evaluating the accuracy
best_logreg = logreg_grid.best_estimator_
test_score = best_logreg.score(X_test, y_test)

print("Test set accuracy:", test_score)
```

Test set accuracy: 1.0

The accuracy is 1.0, very nice.