

20250429_01

April 29, 2025

```
[245]: import pandas as pd
```

```
data = pd.read_csv('data.csv')
```

```
[247]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 29 entries, 0 to 28
Data columns (total 9 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   Country Name    26 non-null    object
 1   Country Code    24 non-null    object
 2   Series Name     24 non-null    object
 3   Series Code     24 non-null    object
 4   2016 [YR2016]   24 non-null    float64
 5   2017 [YR2017]   24 non-null    float64
 6   2018 [YR2018]   24 non-null    float64
 7   2019 [YR2019]   24 non-null    float64
 8   2020 [YR2020]   24 non-null    object
dtypes: float64(4), object(5)
memory usage: 2.2+ KB
```

0.0.1 Structure Overview

- **Columns:** 9 total — Country Name, Country Code, Series Name, Series Code, and data from 2016 to 2020.
- **Rows:** 29 total, but only 24 expected (8 countries \times 3 indicators).
- **Missing Values:**
 - Country Name: 26 non-null \rightarrow 2 additional unexpected values
 - All other key columns have 24 non-null values \rightarrow matches expectation ### Observations
- **Country Name, Country Code, etc:** Correctly typed as object.
- **2016 – 2019:** Correctly typed as float64.
- **2020:** Typed as object, likely due to presence of empty strings or non-numeric text.

```
[250]: data.head(29)
```

[250] :

	Country Name	Country Code	\
0	Argentina	ARG	
1	Argentina	ARG	
2	Argentina	ARG	
3	Brazil	BRA	
4	Brazil	BRA	
5	Brazil	BRA	
6	Chile	CHL	
7	Chile	CHL	
8	Chile	CHL	
9	Paraguay	PRY	
10	Paraguay	PRY	
11	Paraguay	PRY	
12	Uruguay	URY	
13	Uruguay	URY	
14	Uruguay	URY	
15	Bolivia	BOL	
16	Bolivia	BOL	
17	Bolivia	BOL	
18	Peru	PER	
19	Peru	PER	
20	Peru	PER	
21	Ecuador	ECU	
22	Ecuador	ECU	
23	Ecuador	ECU	
24	NaN	NaN	
25	NaN	NaN	
26	NaN	NaN	
27	Data from database: Sustainable Development Go...	NaN	
28	Last Updated: 07/22/2022	NaN	

	Series Name	Series Code	2016 [YR2016]	\
0	Access to electricity (% of population)	EG.ELC.ACCS.ZS	9.984958e+01	
1	CO2 emissions (metric tons per capita)	EN.ATM.CO2E.PC	4.201846e+00	
2	GDP (constant 2015 US\$)	NY.GDP.MKTP.KD	5.823766e+11	
3	Access to electricity (% of population)	EG.ELC.ACCS.ZS	9.970000e+01	
4	CO2 emissions (metric tons per capita)	EN.ATM.CO2E.PC	2.168575e+00	
5	GDP (constant 2015 US\$)	NY.GDP.MKTP.KD	1.743173e+12	
6	Access to electricity (% of population)	EG.ELC.ACCS.ZS	1.000000e+02	
7	CO2 emissions (metric tons per capita)	EN.ATM.CO2E.PC	4.749830e+00	
8	GDP (constant 2015 US\$)	NY.GDP.MKTP.KD	2.467477e+11	
9	Access to electricity (% of population)	EG.ELC.ACCS.ZS	9.840000e+01	
10	CO2 emissions (metric tons per capita)	EN.ATM.CO2E.PC	1.059329e+00	
11	GDP (constant 2015 US\$)	NY.GDP.MKTP.KD	3.775688e+10	
12	Access to electricity (% of population)	EG.ELC.ACCS.ZS	9.970000e+01	
13	CO2 emissions (metric tons per capita)	EN.ATM.CO2E.PC	1.904128e+00	
14	GDP (constant 2015 US\$)	NY.GDP.MKTP.KD	5.417453e+10	

15	Access to electricity (% of population)	EG.ELC.ACCS.ZS	9.180000e+01
16	CO2 emissions (metric tons per capita)	EN.ATM.CO2E.PC	1.995137e+00
17	GDP (constant 2015 US\$)	NY.GDP.MKTP.KD	3.440730e+10
18	Access to electricity (% of population)	EG.ELC.ACCS.ZS	9.420000e+01
19	CO2 emissions (metric tons per capita)	EN.ATM.CO2E.PC	1.838580e+00
20	GDP (constant 2015 US\$)	NY.GDP.MKTP.KD	1.973089e+11
21	Access to electricity (% of population)	EG.ELC.ACCS.ZS	9.870000e+01
22	CO2 emissions (metric tons per capita)	EN.ATM.CO2E.PC	2.414027e+00
23	GDP (constant 2015 US\$)	NY.GDP.MKTP.KD	9.807270e+10
24		NaN	NaN
25		NaN	NaN
26		NaN	NaN
27		NaN	NaN
28		NaN	NaN

	2017 [YR2017]	2018 [YR2018]	2019 [YR2019]	2020 [YR2020]
0	1.000000e+02	9.998958e+01	1.000000e+02	100
1	4.071308e+00	3.975772e+00	3.740650e+00	..
2	5.987909e+11	5.831181e+11	5.713045e+11	514772410744.886
3	9.980000e+01	9.970000e+01	9.980000e+01	100
4	2.196418e+00	2.071855e+00	2.057811e+00	..
5	1.766233e+12	1.797737e+12	1.819683e+12	1749103394213.21
6	9.970000e+01	1.000000e+02	1.000000e+02	100
7	4.714020e+00	4.624338e+00	4.821118e+00	..
8	2.500978e+11	2.600768e+11	2.620808e+11	246412987238.941
9	9.930000e+01	9.960000e+01	9.970000e+01	100
10	1.173720e+00	1.217642e+00	1.165425e+00	..
11	3.957302e+10	4.084104e+10	4.067692e+10	40343452707.5908
12	9.980000e+01	9.980000e+01	9.990000e+01	100
13	1.774987e+00	1.896042e+00	1.874785e+00	..
14	5.505636e+10	5.531948e+10	5.551334e+10	52115108174.7165
15	9.180000e+01	9.280000e+01	9.508000e+01	97.5541229248047
16	2.032547e+00	2.046130e+00	1.940398e+00	..
17	3.585076e+10	3.736496e+10	3.819323e+10	34855949803.3644
18	9.480000e+01	9.520000e+01	9.555136e+01	99.3118133544922
19	1.725909e+00	1.706510e+00	1.745592e+00	..
20	2.022788e+11	2.103080e+11	2.150202e+11	191469666109.311
21	9.920000e+01	9.870000e+01	9.909000e+01	98.8499984741211
22	2.296645e+00	2.349517e+00	2.261470e+00	..
23	1.003954e+11	1.016898e+11	1.017021e+11	93781977159.7812
24	NaN	NaN	NaN	NaN
25	NaN	NaN	NaN	NaN
26	NaN	NaN	NaN	NaN
27	NaN	NaN	NaN	NaN
28	NaN	NaN	NaN	NaN

0.0.2 Observations

1. Country Name and Country Code are redundant — keep one.
2. Series Name and Series Code are also redundant — keep one.
3. As shown in **2020[YR2020]**, e.g., All CO₂ emissions are blank, but there are also float and NaN → **2020** is typed as object
4. Indicators have different scales:
 - Electricity access: max ~100
 - CO₂ emissions: small values (single digit)
 - GDP: large absolute values → **may require standardization**

```
[253]: data.describe()
```

```
[253]:
```

	2016 [YR2016]	2017 [YR2017]	2018 [YR2018]	2019 [YR2019]
count	2.400000e+01	2.400000e+01	2.400000e+01	2.400000e+01
mean	1.247507e+11	1.270115e+11	1.286023e+11	1.293406e+11
std	3.682678e+11	3.736070e+11	3.789742e+11	3.826451e+11
min	1.059329e+00	1.173720e+00	1.217642e+00	1.165425e+00
25%	3.754891e+00	3.627642e+00	3.569208e+00	3.370855e+00
50%	9.920000e+01	9.950000e+01	9.965000e+01	9.975000e+01
75%	4.186130e+10	4.344385e+10	4.446065e+10	4.438602e+10
max	1.743173e+12	1.766233e+12	1.797737e+12	1.819683e+12

0.0.3 Observations

1. `.describe()` not useful due to inconsistent units and missing values.

Next step: filter valid rows, fix 2020 column type, and restructure data.

```
[256]: # First we filter valid rows with Country Name
data['Country Name'].unique()
```

```
[256]: array(['Argentina', 'Brazil', 'Chile', 'Paraguay', 'Uruguay', 'Bolivia',
        'Peru', 'Ecuador', nan,
        'Data from database: Sustainable Development Goals (SDGs)',
        'Last Updated: 07/22/2022'], dtype=object)
```

```
[258]: target_countries = ['Argentina', 'Brazil', 'Chile', 'Paraguay', 'Uruguay',
        ↪ 'Bolivia', 'Peru', 'Ecuador']
data = data[data['Country Name'].isin(target_countries)]
```

```
[260]: # Checking if the cleaning is OK
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 24 entries, 0 to 23
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
#   ...
```

```

---  -----
0   Country Name    24 non-null    object
1   Country Code    24 non-null    object
2   Series Name     24 non-null    object
3   Series Code     24 non-null    object
4   2016 [YR2016]   24 non-null    float64
5   2017 [YR2017]   24 non-null    float64
6   2018 [YR2018]   24 non-null    float64
7   2019 [YR2019]   24 non-null    float64
8   2020 [YR2020]   24 non-null    object
dtypes: float64(4), object(5)
memory usage: 1.9+ KB

```

Looking good, we shall proceed.

```
[263]: data['Series Name'].unique()
```

```
[263]: array(['Access to electricity (% of population)',
            'CO2 emissions (metric tons per capita)',
            'GDP (constant 2015 US$)'], dtype=object)
```

No need to clean, nice. Now we drop some redundant columns.

```
[266]: data = data.drop(columns = ['Country Code', 'Series Code'])
```

```
[268]: # Checking again
data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Index: 24 entries, 0 to 23
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   Country Name    24 non-null    object
1   Series Name     24 non-null    object
2   2016 [YR2016]   24 non-null    float64
3   2017 [YR2017]   24 non-null    float64
4   2018 [YR2018]   24 non-null    float64
5   2019 [YR2019]   24 non-null    float64
6   2020 [YR2020]   24 non-null    object
dtypes: float64(4), object(3)
memory usage: 1.5+ KB

```

Looking good, yipee. Now we clean 2020

```
[271]: data['2020 [YR2020]'] = pd.to_numeric(data['2020 [YR2020]'], errors = 'coerce')
```

```
[273]: # Checking again
data['2020 [YR2020]'].dtype
```

```
[273]: dtype('float64')
```

```
[275]: data['2020 [YR2020]'].isna().sum()
```

```
[275]: 8
```

Now all blanks in 2020 are NaN. I choose to just keep them for now.

Next I make the data easier to read.

```
[298]: # Step 1: Making it tidy
data_melted = data.melt(id_vars = ['Country Name', 'Series Name'],
                        value_vars = ['2016 [YR2016]', '2017 [YR2017]', '2018 [YR2018]', '2019 [YR2019]', '2020 [YR2020]'],
                        var_name = 'year', value_name = 'value')

# Step 2: Rearranging
data_final = data_melted.pivot_table(index = ['Country Name', 'year'],
                                      columns = 'Series Name',
                                      values = 'value').reset_index()
```

```
[300]: # Check
data_final.head()
```

```
[300]: Series Name Country Name      year \
0      Argentina  2016 [YR2016]
1      Argentina  2017 [YR2017]
2      Argentina  2018 [YR2018]
3      Argentina  2019 [YR2019]
4      Argentina  2020 [YR2020]
```

```
Series Name  Access to electricity (% of population) \
0      99.849579
1      100.000000
2      99.989578
3      100.000000
4      100.000000
```

```
Series Name  CO2 emissions (metric tons per capita)  GDP (constant 2015 US$)
0      4.201846      5.823766e+11
1      4.071308      5.987909e+11
2      3.975772      5.831181e+11
3      3.740650      5.713045e+11
4      NaN      5.147724e+11
```

```
[302]: # Exporting the data
data_final.to_csv('cleaned_data.csv', index = False)
```