

20250604_01

June 4, 2025

```
[1]: import pandas as pd
```

```
[4]: # Load data
data = pd.read_csv('melb_data.csv')

# Drop rows with missing values (simple approach for now)
data = data.dropna()
```

```
[19]: # Take a look at the data
data.describe()
```

```
[19]:
```

	Rooms	Price	Distance	Postcode	Bedroom2 \
count	6196.000000	6.196000e+03	6196.000000	6196.000000	6196.000000
mean	2.931407	1.068828e+06	9.751097	3101.947708	2.902034
std	0.971079	6.751564e+05	5.612065	86.421604	0.970055
min	1.000000	1.310000e+05	0.000000	3000.000000	0.000000
25%	2.000000	6.200000e+05	5.900000	3044.000000	2.000000
50%	3.000000	8.800000e+05	9.000000	3081.000000	3.000000
75%	4.000000	1.325000e+06	12.400000	3147.000000	3.000000
max	8.000000	9.000000e+06	47.400000	3977.000000	9.000000

	Bathroom	Car	Landsize	BuildingArea	YearBuilt \
count	6196.000000	6196.000000	6196.000000	6196.000000	6196.000000
mean	1.576340	1.573596	471.006940	141.568645	1964.081988
std	0.711362	0.929947	897.449881	90.834824	38.105673
min	1.000000	0.000000	0.000000	0.000000	1196.000000
25%	1.000000	1.000000	152.000000	91.000000	1940.000000
50%	1.000000	1.000000	373.000000	124.000000	1970.000000
75%	2.000000	2.000000	628.000000	170.000000	2000.000000
max	8.000000	10.000000	37000.000000	3112.000000	2018.000000

	Lattitude	Longtitude	Propertycount
count	6196.000000	6196.000000	6196.000000
mean	-37.807904	144.990201	7435.489509
std	0.075850	0.099165	4337.698917
min	-38.164920	144.542370	389.000000
25%	-37.855438	144.926198	4383.750000
50%	-37.802250	144.995800	6567.000000

75%	-37.758200	145.052700	10175.000000
max	-37.457090	145.526350	21650.000000

```
[12]: # Select features and target
y = data['Price']

features = ['Rooms', 'Bathroom', 'Landsize', 'BuildingArea', 'YearBuilt',
           ↪ 'Latitude', 'Longitude']
X = data[features]
```

```
[10]: # Split the data
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state = 42)
```

```
[14]: # First : linear regression
from sklearn.linear_model import LinearRegression

linear_model = LinearRegression()
linear_model.fit(X_train, y_train)
```

```
[14]: LinearRegression()
```

```
[30]: # Use MAE to evaluate
from sklearn.metrics import mean_absolute_error


linear_predictions = linear_model.predict(X_test)
linear_mae = mean_absolute_error(y_test, linear_predictions)
print(f"MAE: {linear_mae:.0f}")
```

MAE: 293072

```
[31]: # Make plot
import matplotlib.pyplot as plt

# Residuals
linear_residuals = y_test - linear_predictions

plt.figure(figsize = (8, 5))
plt.scatter(linear_predictions, linear_residuals, alpha = 0.5)
plt.axhline(y = 0, color = 'r', linestyle = '--')
plt.xlabel("Predicted Price")
plt.ylabel("Residuals")
plt.title("Residual Plot: Linear Regression")
plt.show()
```



```
[26]: # Second : Decision Tree
from sklearn.tree import DecisionTreeRegressor

tree_model = DecisionTreeRegressor(random_state = 42)
tree_model.fit(X_train, y_train)
```

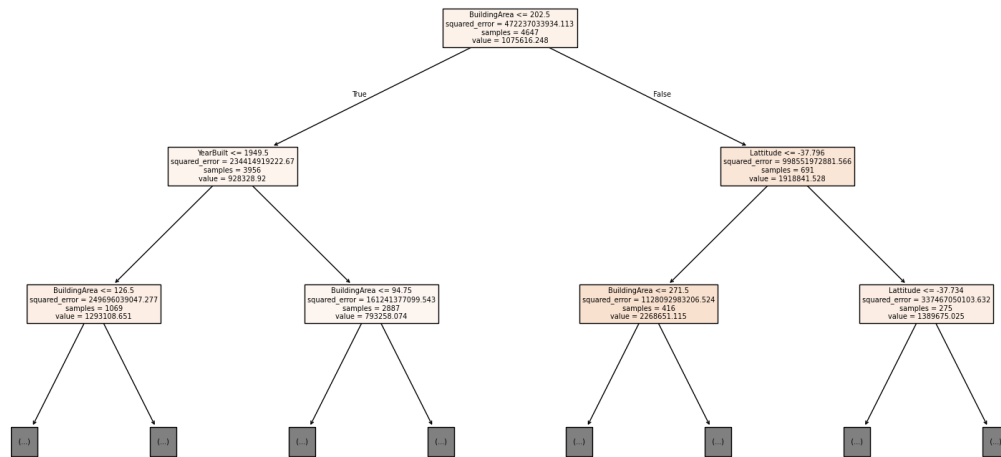
```
[26]: DecisionTreeRegressor(random_state=42)
```

```
[32]: tree_predictions = tree_model.predict(X_test)
tree_mae = mean_absolute_error(y_test, tree_predictions)
print(f"Decision Tree MAE: {tree_mae:.0f}")
```

Decision Tree MAE: 247902

```
[35]: # Make plot
from sklearn.tree import plot_tree

plt.figure(figsize = (20, 10))
plot_tree(tree_model, feature_names = features, filled = True, max_depth = 2)
plt.show()
```



```
[37]: # Try to make it better
pruned_tree = DecisionTreeRegressor(max_depth = 5, random_state = 42)
pruned_tree.fit(X_train, y_train)

pruned_predictions = pruned_tree.predict(X_test)
pruned_mae = mean_absolute_error(y_test, pruned_predictions)
print(f"Pruned Tree MAE: {pruned_mae:.0f}")
```

Pruned Tree MAE: 260739