

20250610_01

June 10, 2025

```
[29]: # Main goal : predict whether a passenger survived or not
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[21]: # Load dataset
data = sns.load_dataset('titanic')
data.head()
```

```
[21]:
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	\
0	0	3	male	22.0	1	0	7.2500	S	Third	
1	1	1	female	38.0	1	0	71.2833	C	First	
2	1	3	female	26.0	0	0	7.9250	S	Third	
3	1	1	female	35.0	1	0	53.1000	S	First	
4	0	3	male	35.0	0	0	8.0500	S	Third	

	who	adult_male	deck	embark_town	alive	alone
0	man	True	NaN	Southampton	no	False
1	woman	False	C	Cherbourg	yes	False
2	woman	False	NaN	Southampton	yes	True
3	woman	False	C	Southampton	yes	False
4	man	True	NaN	Southampton	no	True

```
[22]: # Preview
print(data.shape)
print(data.dtypes)
```

```
(891, 15)
survived          int64
pclass            int64
sex              object
age              float64
sibsp            int64
parch            int64
fare             float64
embarked         object
class            category
who              object
```


```
adult_male      bool
deck            category
embark_town     object
alive           object
alone          bool
dtype: object
```

```
[23]: # Check missing values
data.isnull().sum()
```

```
[23]: survived      0
pclass            0
sex              0
age             177
sibsp            0
parch            0
fare            0
embarked         2
class            0
who              0
adult_male       0
deck            688
embark_town      2
alive            0
alone            0
dtype: int64
```

```
[24]: # Since deck contains too many missing values, I just choose to drop it. Yippee.
# And also embarked and embarked town are the same (I think so ?). So I just
      ↪ decide to leave one.
data = data.drop(columns = ['deck', 'embark_town'])
```

```
[25]: # Take a look at the spread of age.
sns.histplot(data['age'], kde = True)
plt.title('Age Distribution')
plt.show()
```



```
[26]: # Little bit skewed, so fill age with median
data['age'] = data['age'].fillna(data['age'].median())

# Fill embarked with mode, I mean, what else ? (Acutually, what else ? I didn't
↳ learn that much.)
data['embarked'] = data['embarked'].fillna(data['embarked'].mode()[0])
```

```
[28]: # Confirm
data.isnull().sum()
```

```
[28]: survived    0
      pclass      0
      sex         0
      age         0
      sibsp       0
      parch       0
      fare        0
      embarked    0
      class       0
      who         0
```

```
adult_male    0
alive         0
alone         0
dtype: int64
```

```
[31]: data.head()
```

```
[31]:   survived  pclass    sex  age  sibsp  parch    fare embarked  class \
0         0      3   male  22.0     1     0   7.2500          S  Third
1         1      1 female  38.0     1     0  71.2833          C  First
2         1      3 female  26.0     0     0   7.9250          S  Third
3         1      1 female  35.0     1     0  53.1000          S  First
4         0      3   male  35.0     0     0   8.0500          S  Third

      who  adult_male  alive  alone
0   man         True    no  False
1 woman        False   yes  False
2 woman        False   yes   True
3 woman        False   yes  False
4   man         True    no   True
```

```
[32]: # Our main goal is to predict survived
y = data['survived']
```

```
[34]: # pclass and class is the same thing, we choose to keep pclass since is
      ↪numerical.
# also alive is equal to survived, so I shall drop both.
# adult_male can be define via age and sex, so we don't really need it.
# who can also be define via age and sex, drop it too.
drop_cols = ['class', 'survived', 'alive', 'adult_male', 'who']
X = data
X = X.drop(columns = drop_cols)
```

```
[35]: X.head()
```

```
[35]:   pclass    sex  age  sibsp  parch    fare embarked  alone
0      3   male  22.0     1     0   7.2500          S  False
1      1 female  38.0     1     0  71.2833          C  False
2      3 female  26.0     0     0   7.9250          S   True
3      1 female  35.0     1     0  53.1000          S  False
4      3   male  35.0     0     0   8.0500          S   True
```

```
[36]: # Split time
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2,
      ↪random_state = 42)
```

```
[39]: from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier

# Identify column types
# Somehow we can leave Boolean alone because the model is smart enough to
    ↪ handle it?
num_cols = ['pclass', 'age', 'sibsp', 'parch', 'fare']
cat_cols = ['sex', 'embarked']

# Create preprocessor
preprocessor = ColumnTransformer(transformers = [('num', StandardScaler(),
    ↪ num_cols),
                                                ('cat', OneHotEncoder(drop =
    ↪ 'first'), cat_cols)])

# Define pipelines
lr_pipeline = Pipeline(steps = [('preprocessor', preprocessor),
                                ('classifier', LogisticRegression(max_iter=
    ↪ 1000))])

rf_pipeline = Pipeline(steps = [('preprocessor', preprocessor),
                                ('classifier',
    ↪ RandomForestClassifier(random_state = 42))])
```

```
[41]: # Time to evaluate
from sklearn.metrics import accuracy_score, precision_score, recall_score,
    ↪ f1_score, confusion_matrix
```

```
[42]: lr_pipeline.fit(X_train, y_train)
lr_prediction = lr_pipeline.predict(X_test)
```

```
[43]: rf_pipeline.fit(X_train, y_train)
rf_prediction = rf_pipeline.predict(X_test)
```

```
[47]: # Since it's a bit long, 'd better to define a function first
def print_metrics(y_true, y_prediction, model_name):
    print(f"\n{model_name} Performance:")
    print("Accuracy :", accuracy_score(y_true, y_prediction))
    print("Precision:", precision_score(y_true, y_prediction))
    print("Recall   :", recall_score(y_true, y_prediction))
    print("F1 Score :", f1_score(y_true, y_prediction))
```

```
[48]: print_metrics(y_test, lr_prediction, "Logistic Regression")
print_metrics(y_test, rf_prediction, "Random Forest")
```

Logistic Regression Performance:

Accuracy : 0.8100558659217877

Precision: 0.7857142857142857

Recall : 0.7432432432432432

F1 Score : 0.7638888888888888

Random Forest Performance:

Accuracy : 0.8212290502793296

Precision: 0.8

Recall : 0.7567567567567568

F1 Score : 0.7777777777777778

```
[49]: # Making confusion matrix
cm = confusion_matrix(y_test, rf_prediction)

sns.heatmap(cm, annot = True, fmt = 'd', cmap = 'Blues')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Random Forest Confusion Matrix")
plt.show()
```

